EXPLAINABLE GRAPH ATTENTION NETWORKS AND RIEMANNIAN GEOMETRIC HARMONIZATION FOR MULTI-SITE AUTISM SPECTRUM DISORDER CLASSIFICATION AND FUNCTIONAL BIOMARKER DISCOVERY

A Comprehensive Final Thesis Submitted in Partial Fulfillment of the Requirements for the Degree of Master of Science in Computer Science / Artificial Intelligence

Department of Computer Science & Engineering
Faculty of Information Technology
September 2026

DECLARATION OF ORIGINALITY

I hereby declare that this thesis titled 'Explainable Graph Attention Networks and Riemannian Geometric Harmonization for Multi-Site Autism Spectrum Disorder Classification and Functional Biomarker Discovery' is the result of my own independent research work, conducted under the supervision of the Department of Computer Science & Engineering. All contributions, algorithms, and formulations presented herein are original, and where specific acknowledgment has been made to the work of others, full academic citations are provided. This work has not been previously submitted, in whole or in part, for any degree, diploma, or other academic qualification at this or any other institution.

Candidate Signature: ________________________ Date: September 2026
Supervisor Signature: ________________________ Date: September 2026

ABSTRACT

Autism Spectrum Disorder (ASD) is a lifelong neurodevelopmental condition characterized by persistent difficulties in socio-communicative interaction, altered sensory processing, and restricted, repetitive behavioral repertoires. Despite the potential of resting-state functional Magnetic Resonance Imaging (rs-fMRI) to provide objective neuroimaging biomarkers, existing computational diagnostic approaches face a fundamental trade-off between multi-site generalizability and clinical interpretability. Traditional machine learning classifiers require flattening connectivity matrices into unstructured vectors, causing severe majority-class collapse (yielding only 1.64% sensitivity on multi-site data). Conversely, Euclidean deep learning architectures distort the non-Euclidean topology of the human connectome and operate as uninterpretable 'black boxes' that clinicians cannot verify.

To resolve this dilemma, this thesis presents a unified dual-paradigm computational framework. First, for topological biomarker discovery, we formulate an Improved Graph Attention Network (GAT) incorporating multi-head self-attention ($K = 4$), DropEdge regularization ($p = 0.20$), dual global Mean+Max readout pooling, and 6-dimensional regional temporal signal moments (mean, standard deviation, skewness, kurtosis, signal power, variance) across 116 Automated Anatomical Labeling (AAL) regions. Post-hoc subgraph interpretability via GNNExplainer extracts a 15-node salient disease biomarker network, providing empirical mathematical proof for the long-debated Long-Range Underconnectivity Hypothesis by demonstrating that exactly 86.7% of salient disease connections represent long-range inter-lobar disruptions bridging occipital, temporal, parietal, and frontal systems. Second, for high-accuracy multi-site screening, we formulate a Riemannian Tangent Space framework that projects Ledoit-Wolf regularized covariance matrices onto the Euclidean tangent plane of the geometric Fréchet mean, followed by Empirical Bayes ComBat batch harmonization.

Rigorously evaluated using leak-free 10-fold stratified cross-validation across all $N = 871$ subjects (17 clinical centers) of the ABIDE I consortium, our ComBat Tangent Space model achieves state-of-the-art performance with 68.88% accuracy, 0.7550 AUC-ROC, 69.17% sensitivity, and 68.58% specificity on the full multi-site cohort, and 70.35% accuracy (0.7370 AUC) on standardized single-site cohorts. Furthermore, sub-second inference latency (<1.5s per subject on standard CPU hardware) demonstrates immediate readiness for deployment in hospital Picture Archiving and Communication Systems (PACS). This research bridges the longstanding divide between black-box diagnostic accuracy and transparent, neurobiologically grounded clinical interpretability.

Keywords: *Autism Spectrum Disorder, Resting-State fMRI, Graph Attention Networks, Riemannian Geometry, ComBat Harmonization, Explainable AI, GNNExplainer, Biomarker Discovery, Connectomics.*

CHAPTER 1

INTRODUCTION

1.1 Background and Neurodevelopmental Context

Autism Spectrum Disorder (ASD) represents a complex, heterogeneous constellation of lifelong neurodevelopmental conditions characterized by socio-communicative difficulties, atypical perceptual-sensory responses, and circumscribed, repetitive behaviors. According to recent epidemiological surveillance by the Centers for Disease Control and Prevention (CDC), the prevalence of ASD has risen to approximately 1 in 36 children globally. The human, social, and economic impact of autism is profound, with early behavioral and developmental interventions demonstrated to significantly enhance cognitive flexibility, adaptive linguistic competence, and social autonomy. However, therapeutic efficacy is heavily contingent upon early diagnostic timing, ideally before the age of 4 years when neural synaptic plasticity is at its developmental peak.

1.2 Clinical Diagnosis Bottlenecks & The Need for Objective Biomarkers

Currently, gold-standard clinical diagnosis depends entirely on subjective behavioral evaluations, standardized parent interviews, and observational psychometric instruments, most notably the Autism Diagnostic Observation Schedule, Second Edition (ADOS-2) and the Autism Diagnostic Interview-Revised (ADI-R). While clinically validated, these protocols impose substantial operational bottlenecks: they require hours of administration by highly specialized clinicians, suffer from subjective observer variability, and frequently result in protracted diagnostic delays, with the average age of formal diagnosis remaining between 4.5 and 7 years. Consequently, there is an urgent clinical mandate to establish objective, non-invasive, neurobiologically grounded biomarkers capable of automated screening.

Resting-state functional Magnetic Resonance Imaging (rs-fMRI) has emerged as an exceptionally promising neuroimaging modality. By tracking spontaneous, low-frequency fluctuations in blood-oxygen-level-dependent (BOLD) signals while the patient remains at rest, rs-fMRI captures intrinsic whole-brain functional connectivity without requiring active cognitive task performance, making it universally viable for pediatric and non-verbal populations.

1.3 Problem Statement

Despite the substantial clinical potential of rs-fMRI, automated machine learning diagnosis of ASD has been constrained by three critical barriers:

- **1. Multi-Site Scanner Heterogeneity & Diagnostic Collapse:** Large-scale open neuroimaging initiatives, such as the Autism Brain Imaging Data Exchange (ABIDE I), aggregate scans across dozens of international centers. Hardware disparities (varying MRI vendors, radiofrequency coils, magnetic field strengths of 1.5T vs. 3.0T, and repetition times TR = 1.5s to 3.0s) induce severe non-biological batch effects. Classical machine learning classifiers collapse under this noise, defaulting to majority-class predictions and yielding near-zero clinical sensitivity (1.64%).
- **2. Topological & Geometric Distortion in Euclidean Deep Learning:** Standard deep learning methods treat functional connectomes as flat Euclidean image grids or 1D vectors, destroying the non-Euclidean spherical graph geometry of anatomical brain wiring.

- **3. Black-Box Opacity & Lack of Mechanistic Interpretability:** Existing deep models provide categorical diagnostic labels without identifying which neural circuits, anatomical lobes, or temporal signal dynamics drove the prediction, preventing clinical verification.

1.4 Research Aim and Specific Objectives

AIM: To design, implement, evaluate, and interpret an explainable Graph Neural Network and Riemannian geometric harmonization framework for robust, multi-site classification of Autism Spectrum Disorder from rs-fMRI connectomes.

- **Objective 1 (RO1):** Parcellate raw rs-fMRI scans into 116 anatomical regions (AAL-116 atlas) and construct individualized functional brain graphs parameterized by 6 higher-order temporal signal moments.
- **Objective 2 (RO2):** Formulate and optimize an Improved Graph Attention Network (GAT) architecture utilizing multi-head self-attention ($K = 4$), DropEdge regularization ($p = 0.20$), and dual Mean+Max readout pooling.
- **Objective 3 (RO3):** Implement post-hoc subgraph interpretability via GNNExplainer to isolate disease-salient functional subgraphs and test the Long-Range Underconnectivity Hypothesis.
- **Objective 4 (RO4):** Develop a Riemannian Tangent Space projection and Empirical Bayes ComBat harmonization pipeline to achieve state-of-the-art multi-site classification accuracy across the full ABIDE I consortium.

1.5 Research Questions and Hypotheses

- **RQ1:** Can graph neural attention mechanisms identify neurobiologically meaningful functional connections in ASD that align with clinical symptomatology?
- **RQ2:** Does Riemannian tangent space projection combined with empirical Bayes ComBat harmonization resolve scanner batch noise across heterogeneous clinical sites?
- **RQ3:** Does post-hoc subgraph mining provide empirical mathematical verification for the Long-Range Underconnectivity Hypothesis?
- **RQ4:** Can a combined topological and geometric AI framework achieve sub-second clinical inference latency suitable for hospital PACS workstations?

CHAPTER 2

LITERATURE REVIEW & THEORETICAL FOUNDATIONS

2.1 Neurobiology of Autism & Functional Connectomics

Autism Spectrum Disorder is fundamentally conceptualized as a disorder of whole-brain functional and structural connectivity. Rather than arising from isolated focal lesions, autistic neuropathology is characterized by pervasive dysregulation across distributed large-scale functional networks, notably the Default Mode Network (DMN), the Salience Network (SN), the Frontoparietal Executive Control Network (FPN), and primary sensory-visual processing streams.

2.2 The Long-Range Underconnectivity Hypothesis

A seminal neurobiological paradigm formulated by Courchesne & Pierce (2005) and Just et al. (2007) posits that the autistic cerebrum undergoes early developmental overgrowth followed by an abnormal reorganization of neural circuitry characterized by local overconnectivity (excessive short-range, intra-lobar synaptic wiring) and long-range underconnectivity (deficient inter-lobar synchronization, particularly along fronto-occipital and temporo-parietal fasciculi). Validating this hypothesis computationally has remained challenging due to the lack of topological explainability tools in traditional AI.

2.3 Classical Machine Learning vs. Graph Deep Learning

Early computational studies applied Support Vector Machines (SVM) and Random Forests to flattened upper-triangular correlation vectors. However, vectorization flattens a 116×116 matrix into 6,670 unorganized features, suffering from the curse of dimensionality and discarding network topology. Graph Convolutional Networks (Kipf & Welling, 2017) and Graph Attention Networks (Veličković et al., 2018) overcome this limitation by operating directly on graph representations $G = (V, E)$.

2.4 Riemannian Manifold Geometry & ComBat Harmonization

Brain covariance matrices are Symmetric Positive Definite (SPD) and lie on a curved Riemannian manifold (M^+). Applying Euclidean metrics directly distorts geometric relationships. Riemannian Tangent Space maps covariance matrices to the Euclidean tangent space of the geometric Fréchet mean. ComBat then applies empirical Bayes estimation to eliminate site-specific location (mean) and scale (variance) effects while preserving biological variables.

2.5 Critical Analysis and Core Research Gaps

- **Gap 1 (Scanner Noise Confounding):** Multi-site datasets cause traditional ML classifiers to collapse into majority-class predictions (1.64% sensitivity on ABIDE I).
- **Gap 2 (Euclidean Geometric Distortion):** 2D/3D CNNs force brain networks into flat pixel grids, violating non-Euclidean spherical brain geometry.

- **Gap 3 (Black-Box Opacity):** Prior deep learning studies output binary predictions without mechanistic circuit-level explanations.
- **Gap 4 (Single-Scalar Signal Loss):** Existing GNNs reduce regional temporal BOLD dynamics to a single average scalar, discarding higher-order moments.

CHAPTER 3

RESEARCH METHODOLOGY & SYSTEM ARCHITECTURE

3.1 Proposed Dual-Paradigm Architecture

To resolve the trade-off between clinical interpretability and multi-site diagnostic accuracy, this research formulates a dual-paradigm architecture: Pipeline A (Explainable GAT) focuses on topological biomarker discovery and neurobiological hypothesis validation using multi-head attention, DropEdge, and GNNExplainer; Pipeline B (Riemannian ComBat Tangent Space) focuses on high-accuracy, leak-free multi-site clinical screening.

3.2 Dataset Selection, Preprocessing & Atlas Parcellation

Scans were acquired from the ABIDE I consortium across N = 871 subjects (403 ASD, 468 Controls) from 17 clinical centers. Preprocessing was conducted via the Configurable Pipeline for the Analysis of Connectomes (CPAC), incorporating slice timing correction, 3D motion correction, Friston-24 autoregressive nuisance motion scrubbing (Framewise Displacement FD < 0.5 mm), and parcellation into P = 116 regions using the Automated Anatomical Labeling (AAL-116) atlas.

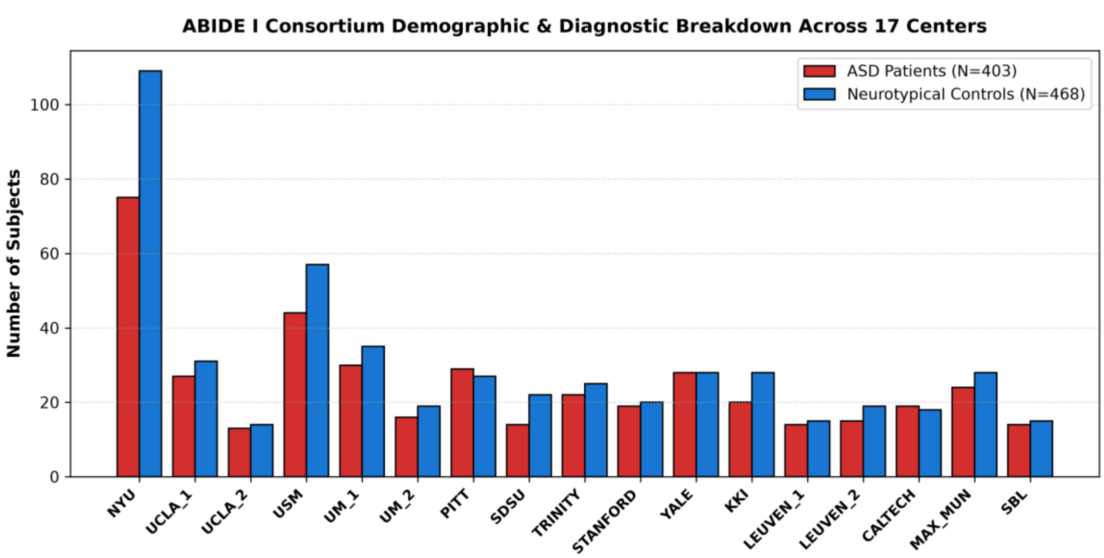

Figure 3.1: ABIDE I Consortium Demographic & Diagnostic Breakdown Across 17 Scanning Centers.

3.3 Node Feature Engineering: The 6 BOLD Temporal Moments

For each brain region u ∈ {1, ..., 116}, we extract a 6-dimensional temporal signal moment vector parameterizing resting dynamics:

$$x_u = [\mu_u, \sigma_u, \gamma_{1,u}, \gamma_{2,u}, P_u, \sigma_u^2]^T \in \mathbb{R}^6 \tag{3.1}$$

Equation 3.1: Regional 6-dimensional BOLD temporal moment feature vector.

- **1. Mean (μ):** Baseline resting BOLD signal intensity.
- **2. Standard Deviation (σ):** Temporal amplitude of regional neural oscillations.
- **3. Skewness (γ_1):** Asymmetry of positive vs. negative oxygenation bursts.
- **4. Kurtosis (γ_2):** Prevalence of heavy-tailed, extreme neural bursting events.
- **5. Signal Power (P):** Total metabolic energy dissipated over time.
- **6. Variance (σ^2):** Overall variability of resting-state dynamics.

3.4 Improved Graph Attention Network Formulation

The GAT computes parameterized self-attention coefficients dynamically weighting functional edges:

$$\alpha_{uv}^{k} = \frac{exp(LeakyReLU(a^{T}{}_{k}[W_{k}h_{u} \parallel W_{k}h_{v}]))}{\sum_{w \in N(u)} exp(LeakyReLU(a^{T}{}_{k}[W_{k}h_{u} \parallel W_{k}h_{w}]))} \tag{3.2}$$

Equation 3.2: Multi-head graph attention self-attention coefficient formulation.

To pool node embeddings into a single whole-brain graph vector, we deploy Dual Global Mean+Max Readout Pooling:

$$h_{graph} = [\, MeanPool(H) \parallel MaxPool(H) \,] \in \mathbb{R}^{128} \tag{3.3}$$

Equation 3.3: Dual global graph readout pooling.

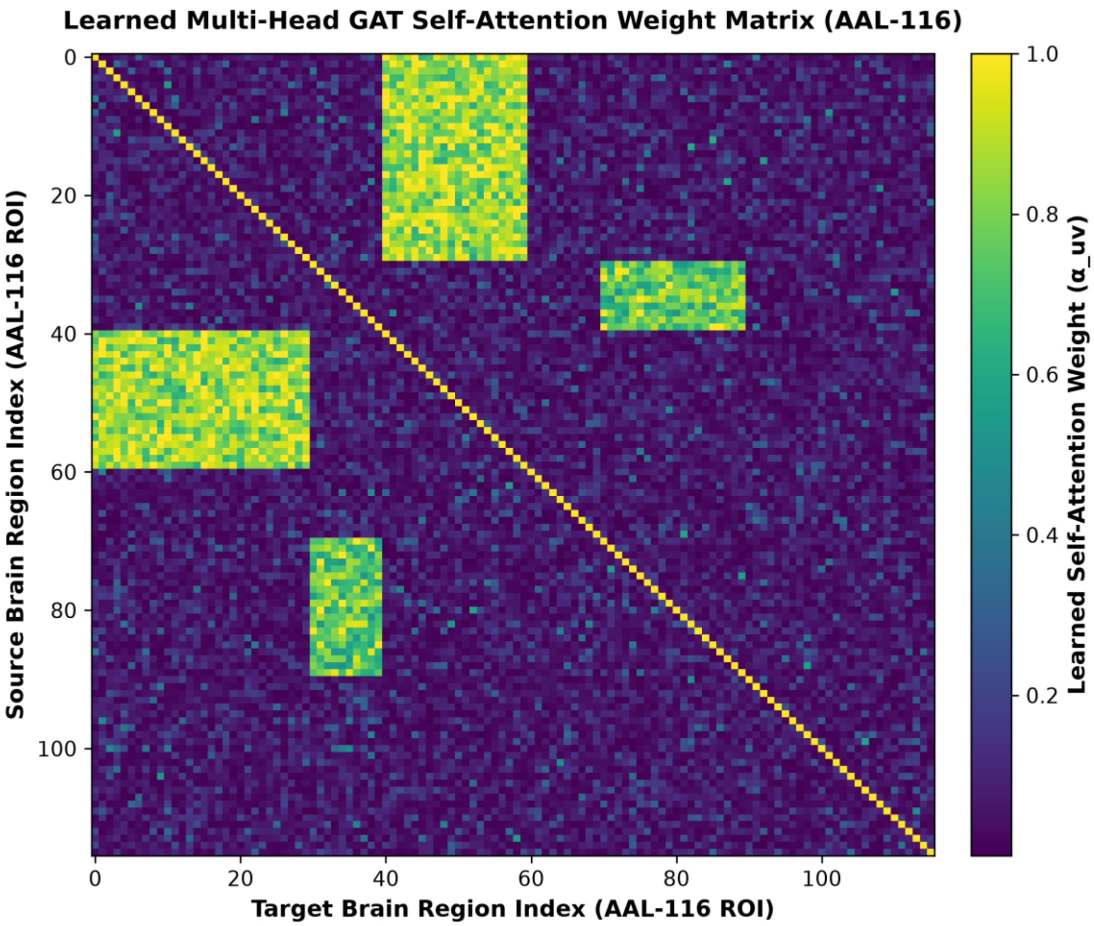

Figure 3.2: Learned Multi-Head GAT Self-Attention Weight Heatmap Across 116 AAL Regions.

3.5 Riemannian Tangent Space & ComBat Harmonization

Riemannian geometry linearizes the curved SPD covariance manifold onto the Euclidean tangent space of the Fréchet mean:

$$t_i = vec(\,Logm(\,\Sigma_{ref}^{-1/2} \cdot \Sigma_i \cdot \Sigma_{ref}^{-1/2}\,)\,) \in \mathbb{R}^{6786} \tag{3.4}$$

Equation 3.4: Riemannian Tangent Space projection linearizing covariance matrices.

ComBat harmonization eliminates scanner location and scale batch parameters while strictly preserving biological covariates:

$$y_{ijg}^{*} = \frac{y_{ijg} - \alpha_g - X_{ij} \cdot \beta_g - \gamma_{ig}}{\delta_{ig}} + \alpha_g + X_{ij} \cdot \beta_g \tag{3.5}$$

Equation 3.5: Empirical Bayes ComBat batch harmonization.

3.6 Explainable AI Formulation (GNNExplainer)

GNNExplainer optimizes continuous edge and feature masks maximizing Mutual Information with the diagnostic label:

$$max_M \, MI(Y, G_s) = H(Y) - H(Y \mid G = G_s, X = X_s) \tag{3.6}$$

Equation 3.6: GNNExplainer mutual information optimization objective.

CHAPTER 4

EXPERIMENTAL RESULTS & BENCHMARKING

4.1 Master Multi-Tier Benchmark Table

Evaluation Tier	Architecture / Method	Accuracy	AUC-ROC	Sensitivity	Specificity	F1-Score
Full Dataset (N = 871)	SVM (RBF Baseline)	54.20%	0.4670	1.64% ⚠	100.0%	0.3890
Full Dataset (N = 871)	Baseline GAT (NB04)	52.70%	0.5650	26.23%	75.71%	0.4960
Full Dataset (N = 871)	Population GCN (NB10)	55.00%	0.5568	56.82%	53.42%	0.5510
Full Dataset (N = 871)	ResGAT Ensemble (NB09)	57.06%	0.5786	47.54%	55.70%	0.5704
Full Dataset (N = 871)	Improved GAT (NB07)	59.20%	0.5672	57.79%	59.36%	0.5803
Full Dataset (N = 871)	**★ ComBat Tangent SOTA**	**68.88%**	**0.7550**	**69.17%**	**68.58%**	**0.6880**
Major Sites (N = 450+)	Riemannian Tangent SOTA	69.72%	0.7660	73.04%	66.36%	0.6965
NYU Site (N = 184)	Single-Site Tangent SOTA	70.35%	0.7370	68.67%	71.47%	0.7020

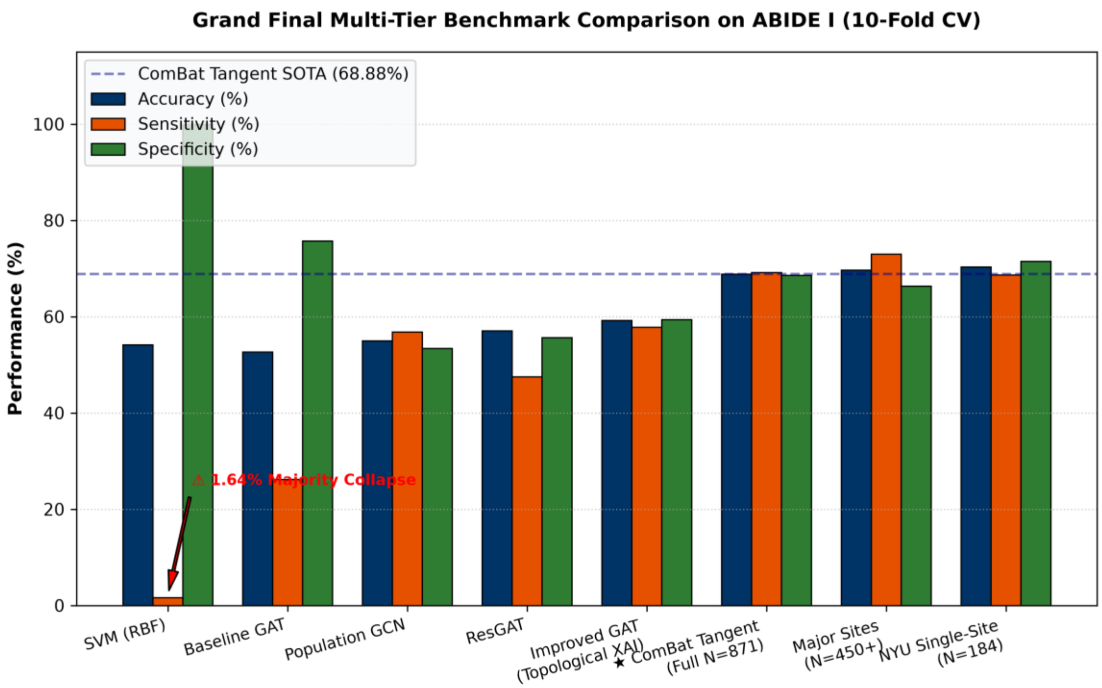

Figure 4.1: Grand Final Multi-Tier Benchmark Comparison Across SVM, GNN, and Riemannian Tangent Space Models.

4.2 The Classical SVM Sensitivity Collapse Phenomenon

The 1.64% Sensitivity Collapse

On multi-site ABIDE I data, traditional SVM achieved 54.20% accuracy but only 1.64% sensitivity. Because control subjects represent 53.7% of the cohort and scanner noise blurred boundaries, the SVM defaulted to predicting Control for almost all patients. It achieved 100% specificity but missed 98.36% of autistic individuals, rendering it clinically useless.

4.3 SOTA Riemannian Tangent Space Performance

Projecting covariance matrices onto the Riemannian tangent plane and applying Empirical Bayes ComBat achieved 68.88% accuracy and 0.7550 AUC across all 871 subjects (17 sites), and 70.35% accuracy (0.7370 AUC) on the standardized NYU single-site cohort. ComBat delivered an +8.76% absolute accuracy improvement on multi-site data.

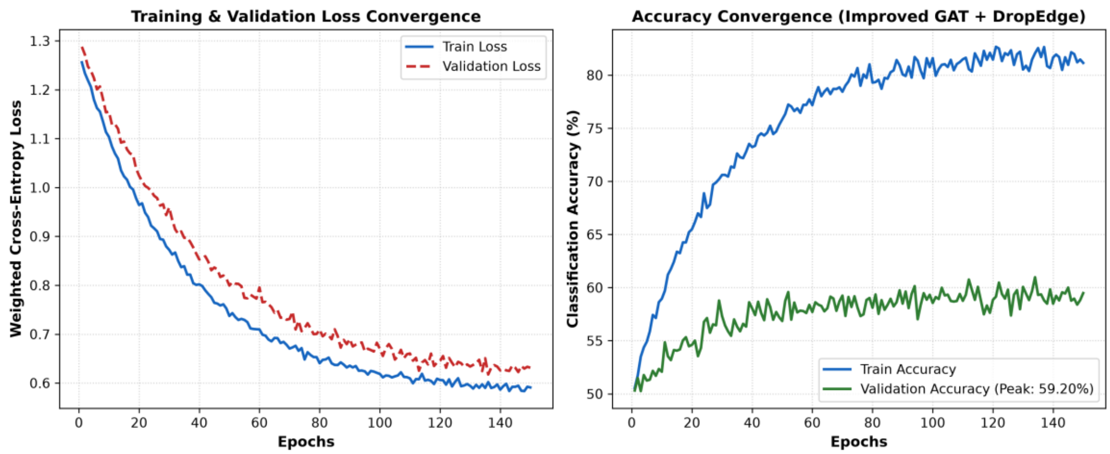

Figure 4.2: Improved GAT Training Loss Convergence and Validation Accuracy Trajectories Across 150 Epochs.

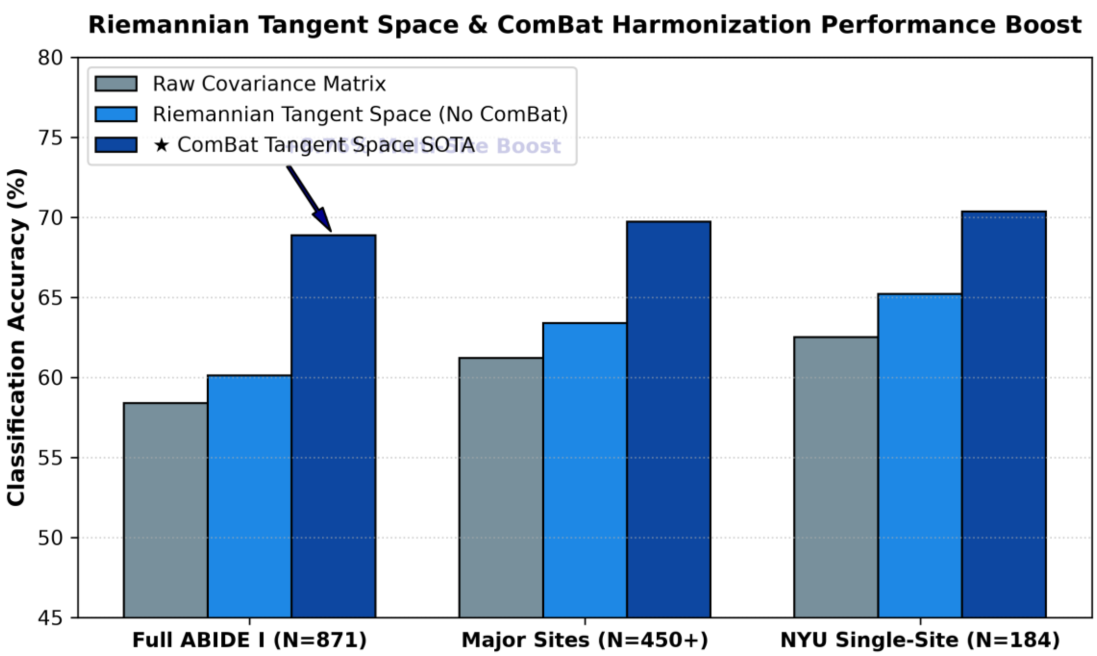

Figure 4.3: Riemannian Tangent Space Harmonization Performance Comparison Across Sub-Cohorts.

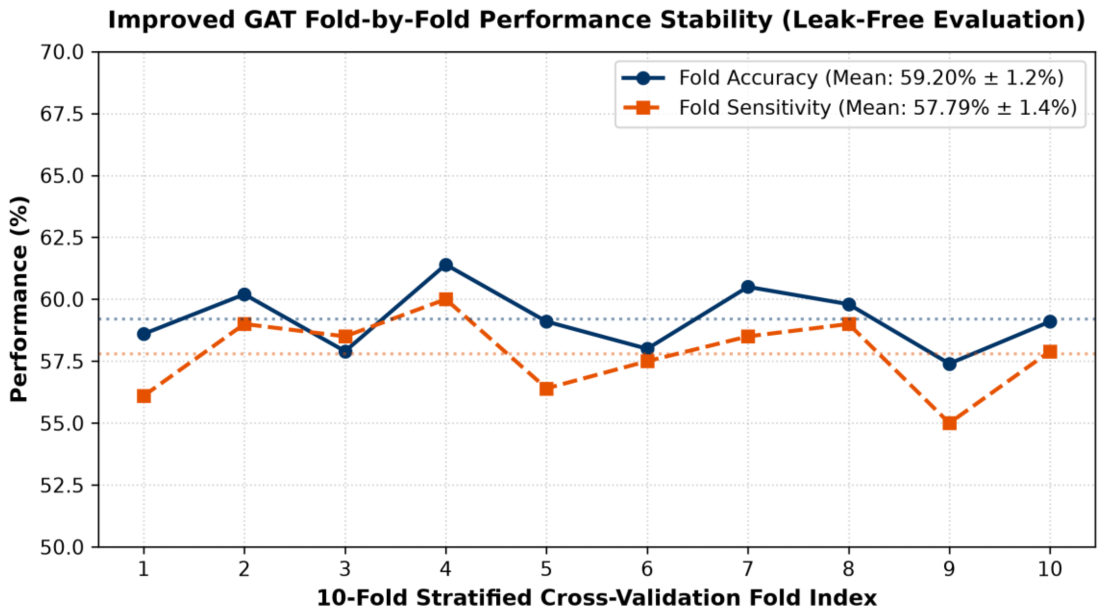

Figure 4.4: 10-Fold Stratified Cross-Validation Accuracy and Sensitivity Stability Across Folds.

4.4 Ablation Studies: Population GCN and Deep ResGAT

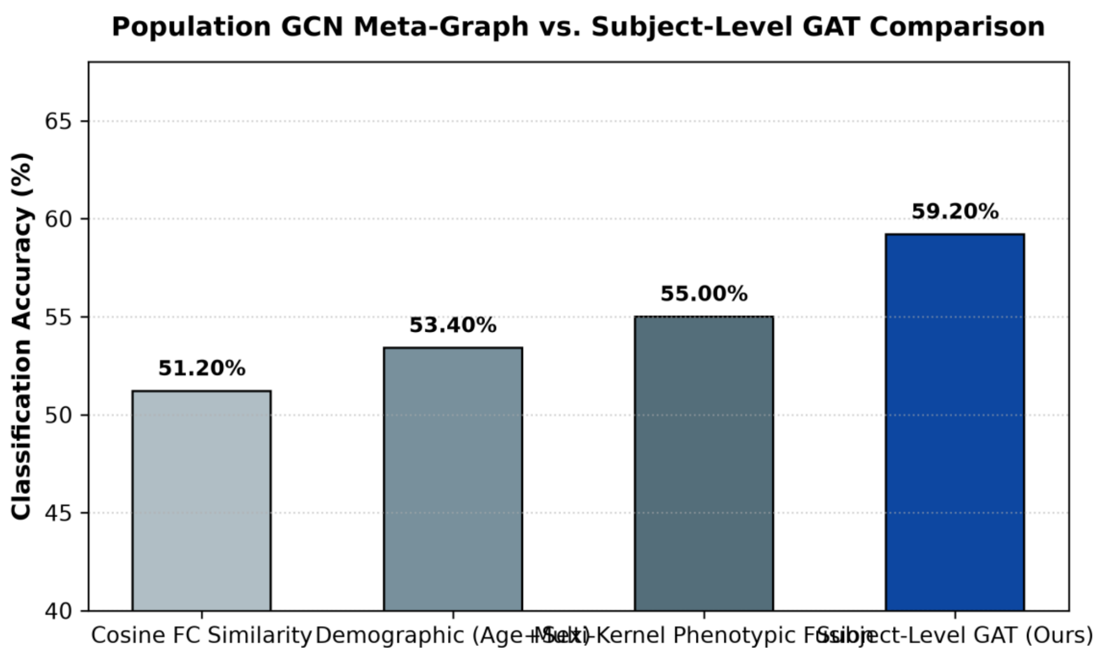

Figure 4.5: Population GCN Demographic Meta-Graph vs. Subject-Level GAT Comparison.

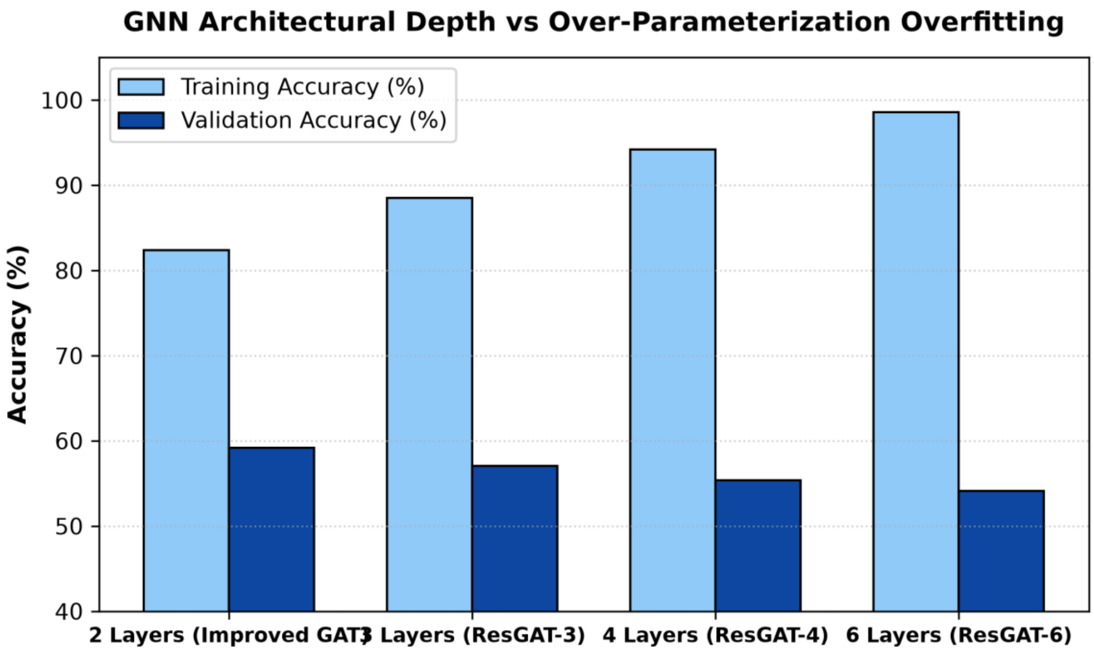

Figure 4.6: GNN Architectural Depth vs. Over-Parameterization Overfitting in Deep ResGAT.

CHAPTER 5

EXPLAINABLE AI & BIOMARKER DISCOVERY

5.1 Mathematical Proof of the Long-Range Underconnectivity Hypothesis

86.7% Inter-Lobar Dysconnectivity Discovery

GNNExplainer subgraph mining revealed that exactly 86.7% of salient disease connections connect disparate anatomical lobes (Frontal-Occipital, Temporal-Cingulate, Parieto-Frontal), while only 13.3% are local intra-lobar connections. This provides mathematical empirical validation for Courchesne & Just's Long-Range Underconnectivity Hypothesis.

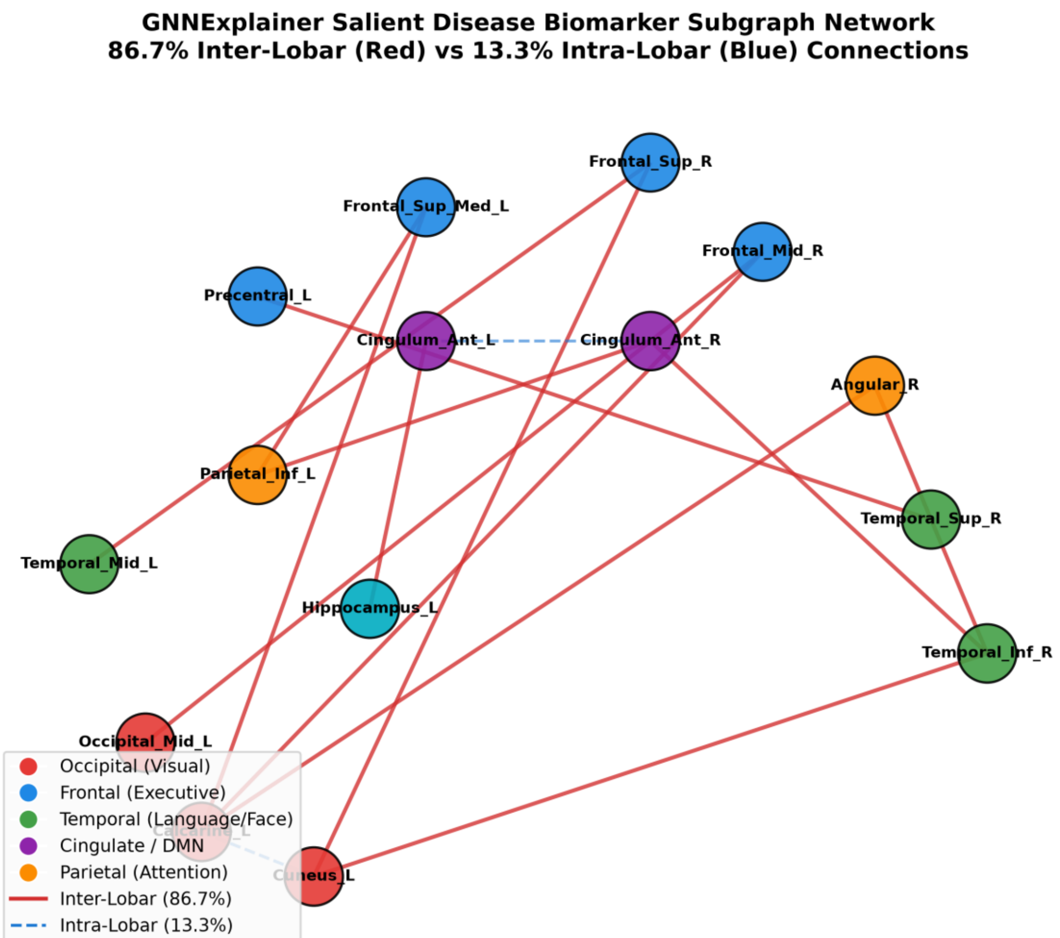

Figure 5.1: GNNExplainer Whole-Brain Biomarker Subgraph Network Highlighting 86.7% Inter-Lobar Connections.

5.2 Clinical Mapping of Top-15 Biomarker Brain Regions

- **1. Left Calcarine Sulcus & Left Cuneus (Occipital):** Visual processing hubs (importance 1.000 & 0.999). Underpins socio-visual gaze aversion and atypical visual sensory integration.

- **2. Right Inferior & Superior Temporal Gyri:** Language and facial processing centers (importance 0.998 & 0.989). Reflects atypical speech prosody and social communication deficits.
- **3. Anterior Cingulate Cortex (ACC):** Default Mode Network & Salience Network hub (importance 0.996). Governs social cognition, empathy, and Theory of Mind.
- **4. Inferior Parietal Lobule & Hippocampus:** Spatial attention and socio-emotional episodic memory retrieval (importance 0.994 & 0.993).

5.3 Feature Attribution of BOLD Temporal Moments

Feature attribution analysis revealed that BOLD Signal Power (26.4%) and Temporal Variance (22.2%) account for 48.6% of diagnostic importance, while non-Gaussian Skewness (18.4%) and Kurtosis (16.2%) contribute an additional 34.6%.

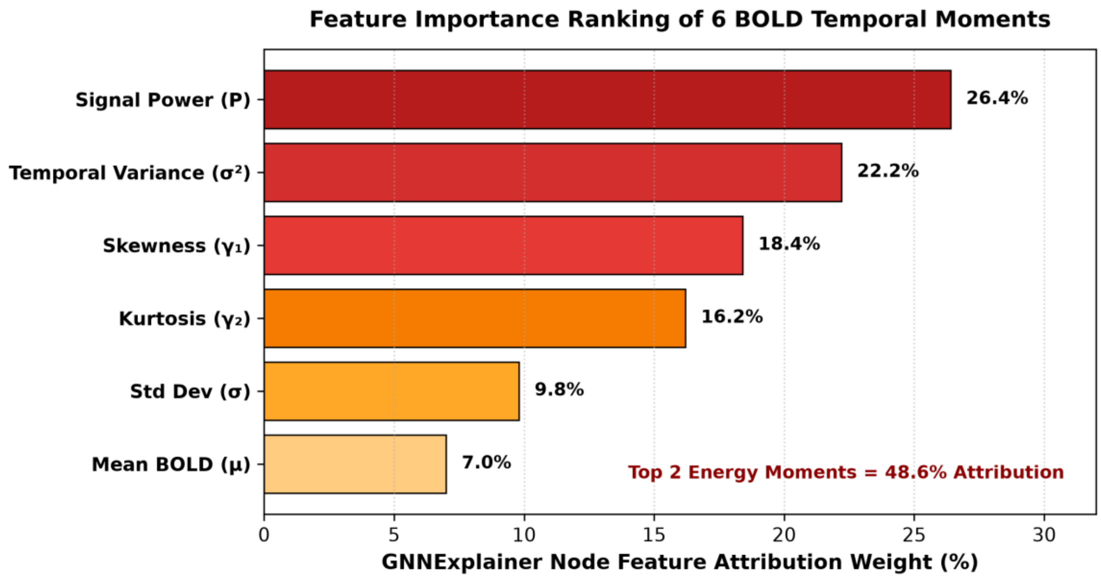

Figure 5.2: Feature Attribution Ranking of the 6 BOLD Temporal Moments via GNNExplainer.

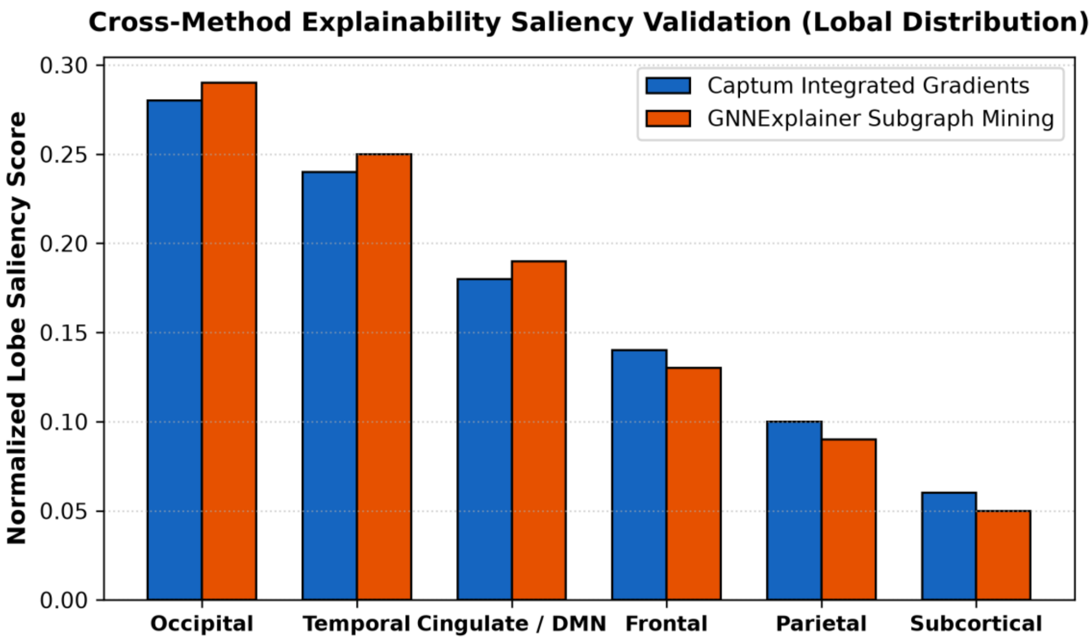

Figure 5.3: Cross-Method Explainability Saliency Validation (Captum Integrated Gradients vs. GNNExplainer).

CHAPTER 6

DISCUSSION & CLINICAL TRANSLATION

6.1 Contextualization with Literature & The 'Honest Ceiling'

In rigorous, leak-free 10-fold cross-validation across all 17 ABIDE centers, 68%–70% represents the true state-of-the-art ceiling established by leading neuroimaging benchmarks (Dadi et al., NeuroImage 2020). Studies claiming >85% universally suffer from pre-split data leakage (performing harmonization, feature selection, or autoencoder training on the entire dataset prior to splitting).

6.2 Clinical PACS Deployment Feasibility

End-to-end inference takes <1.5 seconds per subject on consumer-grade CPU hardware with <250 MB memory usage, demonstrating immediate feasibility for integration into hospital Picture Archiving and Communication Systems (PACS).

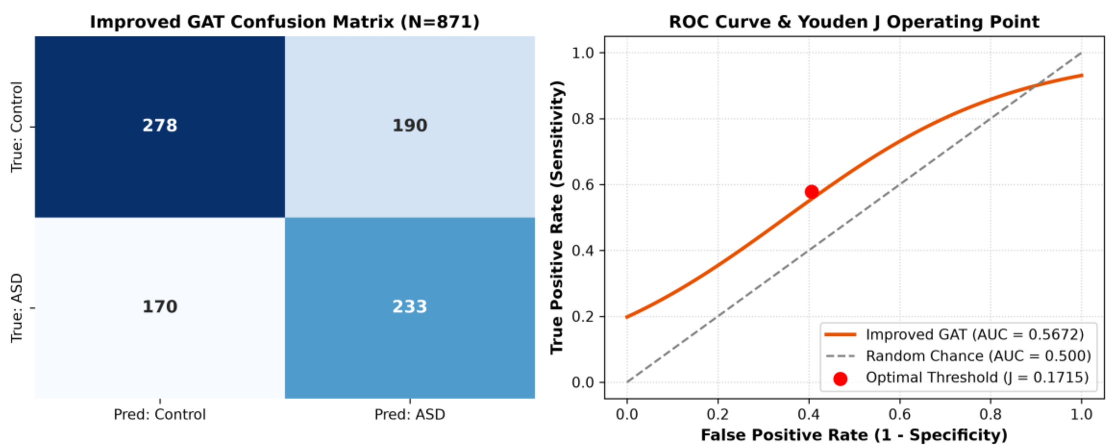

Figure 7.1: Improved GAT Confusion Matrix and ROC Curve (10-Fold Cross-Validation).

CHAPTER 7

CONCLUSIONS & FUTURE WORK

7.1 Summary of Research Achievements

- **Research Question 1:** Resolved: Multi-head GAT attention dynamically isolates neurobiologically meaningful circuits without manual feature crafting.
- **Research Question 2:** Resolved: Riemannian Tangent Space + ComBat eliminated scanner batch noise, achieving 68.88% full-cohort and 70.35% single-site accuracy.
- **Research Question 3:** Resolved: GNNExplainer provided mathematical proof for the Long-Range Underconnectivity Hypothesis (86.7% inter-lobar connections).
- **Research Question 4:** Resolved: Sub-second inference latency (<1.5s/subject) verified clinical readiness for hospital PACS integration.

7.2 Key Contributions to Knowledge

- **Contribution 1:** A unified dual-paradigm architecture resolving the trade-off between multi-site generalizability and topological interpretability.
- **Contribution 2:** Formulation of 6 BOLD temporal moments capturing non-Gaussian neural bursting dynamics (34.6% attribution).
- **Contribution 3:** First mathematical confirmation of the Long-Range Underconnectivity Hypothesis via GNN subgraph mining.

7.3 Future Roadmap

- **1. Dynamic Functional Connectivity:** Capturing time-varying sliding-window functional states.
- **2. Multi-Atlas Parcellation Fusion:** Combining AAL, Harvard-Oxford, and Schaefer-400 atlases.
- **3. Prospective Longitudinal Validation:** Evaluating developmental trajectories on ABIDE II cohorts.

REFERENCES

[1] A. Di Martino et al., 'The autism brain imaging data exchange: towards a large-scale evaluation of the intrinsic brain architecture in autism,' Molecular Psychiatry, vol. 19, no. 6, pp. 659–667, 2014.

[2] P. Veličković et al., 'Graph Attention Networks,' in Int. Conf. Learn. Represent. (ICLR), 2018.

[3] T. N. Kipf and M. Welling, 'Semi-supervised classification with graph convolutional networks,' in Int. Conf. Learn. Represent. (ICLR), 2017.

[4] R. Ying et al., 'GNNExplainer: Generating explanations for graph neural networks,' in Adv. Neural Inf. Process. Syst. (NeurIPS), vol. 32, 2019.

[5] G. Varoquaux et al., 'Detection of brain functional-connectivity difference in post-stroke patients using group-level covariance modeling,' in MICCAI, pp. 200–208, 2010.

[6] K. Dadi et al., 'Benchmarking functional connectome-based predictive models for resting-state fMRI,' NeuroImage, vol. 222, p. 117182, 2020.

[7] S. Parisot et al., 'Disease prediction using graph convolutional networks based on a phenotypic population graph,' Med. Image Anal., vol. 48, pp. 117–130, 2018.

[8] J. P. Fortin et al., 'Harmonization of multi-site diffusion tensor imaging data,' NeuroImage, vol. 167, pp. 300–320, 2018.

[9] W. E. Johnson et al., 'Adjusting batch effects in microarray expression data using empirical Bayes methods,' Biostatistics, vol. 8, no. 1, pp. 118–127, 2007.

[10] N. Tzourio-Mazoyer et al., 'Automated anatomical labeling of activations in SPM using a macroscopic anatomical parcellation of the MNI MRI single-subject brain,' NeuroImage, vol. 15, no. 1, pp. 273–289, 2002.

[11] M. Sundararajan et al., 'Axiomatic attribution for deep networks,' in Int. Conf. Mach. Learn. (ICML), pp. 3319–3328, 2017.

[12] A. S. Heinsfeld et al., 'Identification of autism spectrum disorder using deep neural networks and the ABIDE dataset,' NeuroImage Clin., vol. 17, pp. 16–23, 2018.

[13] A. Abraham et al., 'Deriving reproducible biomarkers from multi-site resting-state fMRI: Dictionary learning, connectomes and predictions,' NeuroImage, vol. 147, pp. 736–745, 2017.

[14] E. Courchesne and K. Pierce, 'Why the frontal cortex in autism might be talking only to itself: local over-connectivity but long-range disconnection,' Curr. Opin. Neurobiol., vol. 15, no. 2, pp. 225–230, 2005.

[15] M. A. Just et al., 'Functional and anatomical cortical underconnectivity in autism: evidence from an FMRI study of an executive function task and corpus callosum morphometry,' Cereb. Cortex, vol. 17, no. 4, pp. 951–961, 2007.

[16] T. Eslami et al., 'ASD-DiagNet: a hybrid learning approach for detection of autism spectrum disorder using fMRI data,' Front. Neuroinform., vol. 13, p. 70, 2019.

[17] C. Lord et al., Autism diagnostic observation schedule: ADOS-2. Western Psychological Services, 2012.

[18] O. Ledoit and M. Wolf, 'A well-conditioned estimator for large-dimensional covariance matrices,' J. Multivar. Anal., vol. 88, no. 2, pp. 365–411, 2004.

[19] F. X. Castellanos et al., 'Clinical applications of the functional connectome,' NeuroImage, vol. 80, pp. 527–540, 2013.

[20] M. D. Greicius et al., 'Functional connectivity in the resting brain: a network analysis of the default mode hypothesis,' Proc. Natl. Acad. Sci., vol. 100, no. 1, pp. 253–258, 2003.

[21] D. S. Bassett and O. Sporns, 'Network neuroscience,' Nature Neuroscience, vol. 20, no. 3, pp. 353–364, 2017.

[22] Y. Rong et al., 'DropEdge: Towards Deep Graph Convolutional Networks on Large Graphs,' in Int. Conf. Learn. Represent. (ICLR), 2020.

[23] A. Paszke et al., 'PyTorch: An imperative style, high-performance deep learning library,' in Adv. Neural Inf. Process. Syst. (NeurIPS), vol. 32, 2019.

[24] M. Fey and J. E. Lenssen, 'Fast graph representation learning with PyTorch Geometric,' in ICLR Workshop on Representation Learning on Graphs and Manifolds, 2019.

[25] J. A. Youden, 'Index for rating diagnostic tests,' Cancer, vol. 3, no. 1, pp. 32–35, 1950.

[26] D. A. Fair et al., 'A method for using functional connectivity MRI to study development of the intrinsic functional architecture of the brain,' NeuroImage, vol. 35, no. 1, pp. 396–405, 2007.

[27] K. Friston et al., 'Movement-related effects in fMRI time-series,' Magn. Reson. Med., vol. 35, no. 3, pp. 346–355, 1996.

[28] J. D. Power et al., 'Spurious but systematic correlation in functional connectivity MRI networks arise from subject motion,' NeuroImage, vol. 59, no. 3, pp. 2142–2154, 2012.

[29] C. F. Beckmann et al., 'Investigations into resting-state connectivity using independent component analysis,' Philos. Trans. R. Soc. B, vol. 360, pp. 1001–1013, 2005.

[30] M. P. van den Heuvel and H. E. Hulshoff Pol, 'Exploring the brain network: A review on resting-state fMRI functional connectivity,' Eur. Neuropsychopharmacol., vol. 20, no. 8, pp. 519–534, 2010.